











SQL Server Administration

Part10



Mohammad Reza Gerami mrgerami@aut.ac.ir gerami@virasec.ir





Extended Events

Extended Events

Extended events is a lightweight performance monitoring system that enables users to collect data needed to monitor and troubleshoot problems in SQL Server. See Extended events overview to learn more about the extended events architecture. This article aims to help the SQL developer who is new to extended events, and who wants create an event session in just a few minutes. By using extended events, you can see details about the inner operations of the SQL system and your application. When you create an extended event session, you tell the system:

Which occurrences you are interested in.

How you want the system to report the data to you







Database Snapshots

Security Solition

A database snapshot is a read-only, static view of a SQL Server database (the source database). The database snapshot is transactionally consistent with the source database as of the moment of the snapshot's creation. A database snapshot always resides on the same server instance as its source database.

BCP



bcp "Person.Address" out d:\Address.txt -S "127.0.0.1" -d AdventureWorks2019 -Usa -Psql -c -T

bcp "Person.Address" in d:\Address.txt -S "127.0.0.1" -d AdventureWorks2019 -Usa -Psql -c -T

https://docs.microsoft.com/en-us/sql/tools/bcp-utility?view=sql-server-ver15



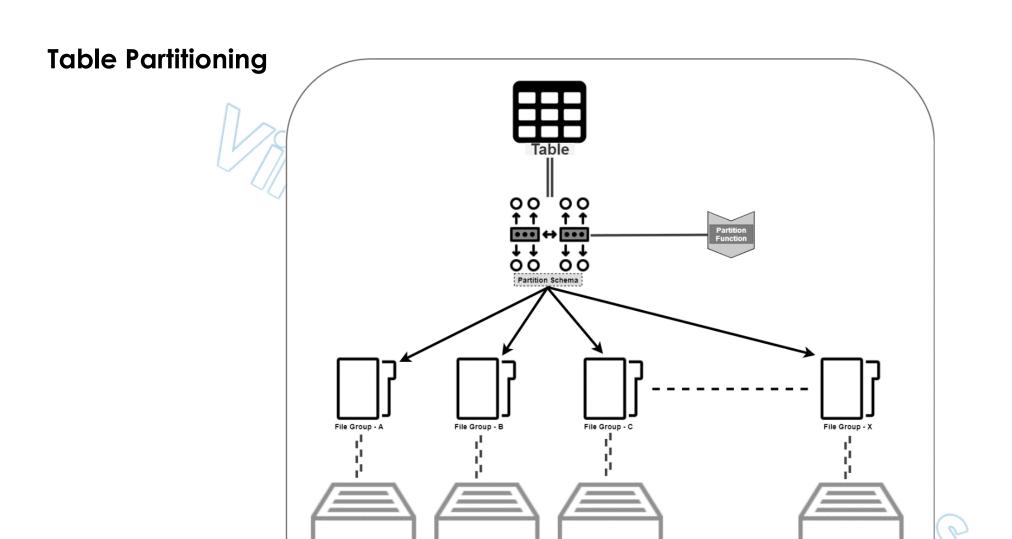
Table Partitioning in SQL Server

Table Partitioning in SQL Server



What is a database table partitioning?

Partitioning is the database process where very large tables are divided into multiple smaller parts. By splitting a large table into smaller, individual tables, queries that access only a fraction of the data can run faster because there is less data to scan. The main of goal of partitioning is to aid in maintenance of large tables and to reduce the overall response time to read and load data for particular SQL operations.





Partitions

Partitions

Partitions

Partitions



5, 5, 5, 7, 6	ROW1 = 2005	ROW1 = 2005	ROW1 = 2007	ROW1 = 2006	ROW1 = 2006	ROW1 = 2005	ROW1 = 2007	ROW1 = 2005	ROW1 = 2015
---------------	-------------	-------------	-------------	-------------	-------------	-------------	-------------	-------------	-------------

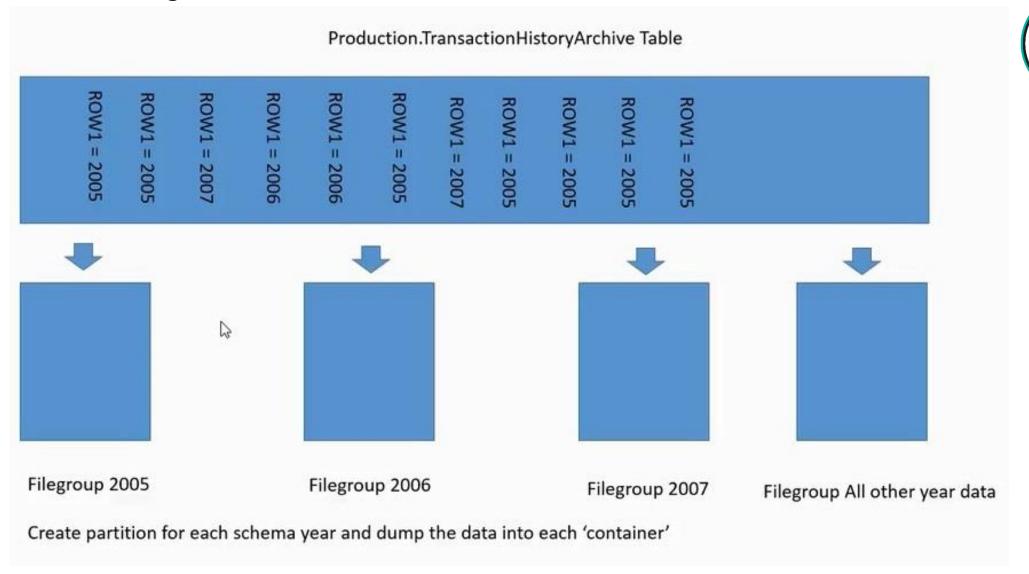


Production.TransactionHistoryArchive Table

ROW1 = 2005	ROW1 = 2005	ROW1 = 2007	ROW1 = 2006	ROW1 = 2006	ROW1 = 2005	ROW1 = 2007	ROW1 = 2005									
G	5	7	6	6	5	Ž	5	5	Ŭ,	ŭ	5	5	5	5	5	

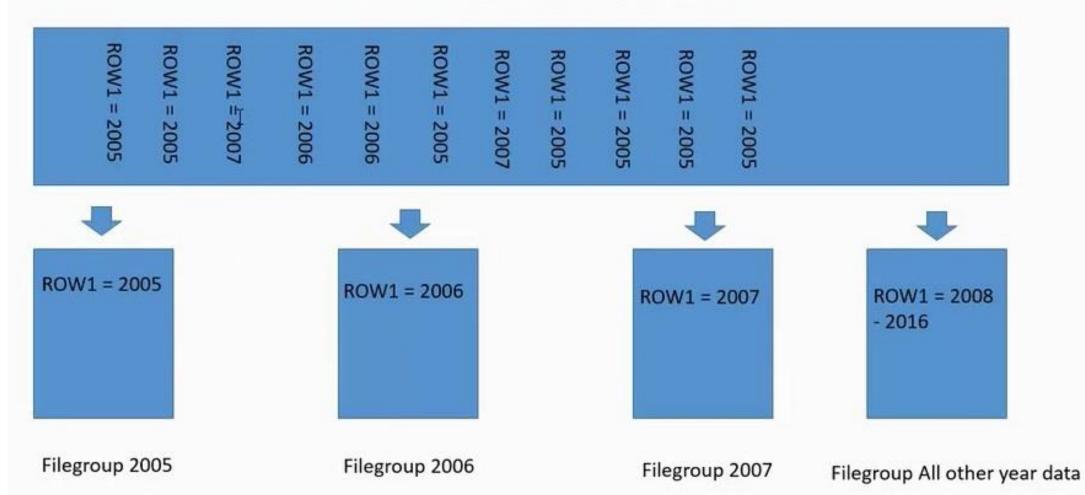
As the table gets bigger and bigger in time, it becomes more difficult to manage; the space required, the speed of data retrieval, the increased I/O activity, the hardware resources, the deletes, the inserts, the updates, the selects, the backups and restores all start to press upon the SQL Server and hardware. To remedy this one issue we can consider table partitioning. Table partitioning, simply put, divides the content of the table in to smaller manageable chunks of data in a file group.







Production.TransactionHistoryArchive Table



Production.TransactionHistoryArchive Table ROW1 = 2005 ROW1 = 2005 ROW1 = 2007 ROW1 = 2006 ROW1 = 2007 ROW1 = 2005 ROW1 = 2005 ROW1 = 2005 ROW1 = 2005 ROW1 = 2008ROW1 = 2006 - 2016 ROW1 = 2007 Archive Filegroup 2005,2006,2007 Partition All other year data Archive all data into a single partition for read only



What is Table Partitioning in SQL Server? Why is partitioning required in SQL Server?

Vital Security Soldier

Table Partition is the logical division of information with the physical distribution in the filegroups.

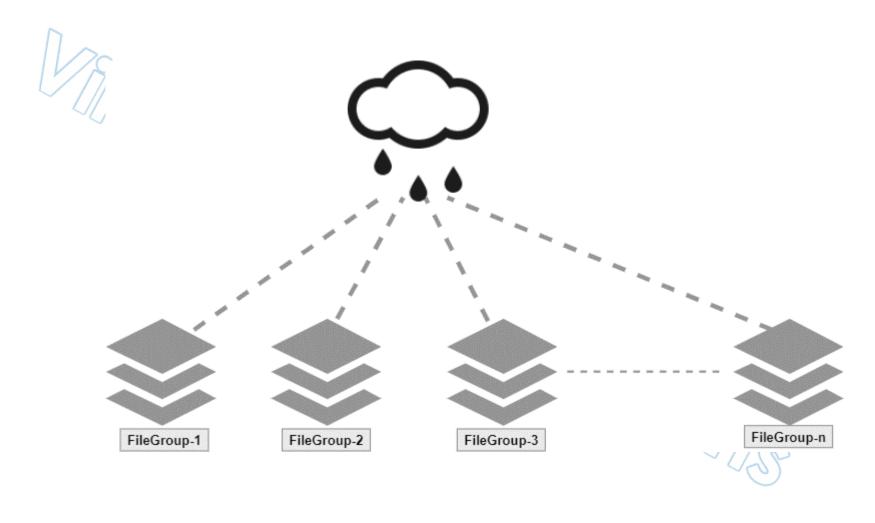
A table can be partitioned by applying the partition schema over the table schema. The role of the partition function is to divide the information in a logical division by partition range and partition scheme indexing the partition range to the filegroup.

After applying the partition on the table, rows of the table will be distributed in different secondary files by the filegroup.



Partition is one of the beneficial approaches for query performance over the large table. Table Index will be a part of each partition in SQL Server to return a quick query response. Actual partition performance will be achieved when you use a query with the partition column because the partition column will target those partitions only, which are required by the partition column. It wouldn't scan the information in all filegroups.





Creating a sample database to perform Table Partitioning in SQL Server:

CREATE DATABASE AutoPartition



ALTER DATABASE AutoPartition ADD FILEGROUP FG_01_2020 GO





Adding Files to each File Group



```
ALTER DATABASE AutoPartition
ADD FILE
(

NAME = [File_012020],
FILENAME = 'C:\Program Files\Microsoft SQL Server\MSSQL13.JRAIYANI\MSSQL\DATA\File_012020.ndf',
    SIZE = 5 MB,
    MAXSIZE = UNLIMITED,
    FILEGROWTH = 10 MB
) TO FILEGROUP FG_01_2020
GO
```

Here, adding one file to each filegroup for data distribution to the physical storage to perform the table partitioning in SQL Server. If you have crores of row per day, then you can add multiple files to the single file group as well, and data will be distributed to the number of files in the filegroup.

Adding a Partition Function with Month wise range

```
USE AutoPartition
GO
CREATE PARTITION FUNCTION [PF_MonthlyPartition] (DATETIME)
AS RANGE RIGHT FOR VALUES
(
'2020-01-31 23:59:59.997', '2020-02-29 23:59:59.997', '2020-03-31 23:59:59.997',
'2020-04-30 23:59:59.997', '2020-05-31 23:59:59.997', '2020-06-30 23:59:59.997',
'2020-07-31 23:59:59.997', '2020-08-31 23:59:59.997', '2020-09-30 23:59:59.997',
'2020-10-31 23:59:59.997', '2020-11-30 23:59:59.997', '2020-12-31 23:59:59.997'
);
```

Here, 12 ranges are defined with the last day of the month and last ms of the day. Users can use the month and year combination as well to perform the table partitioning in SQL Server. But I would recommend to define it with full datetime to perform insert operations quickly. In the above sample code, the partition function is defined with the Right direction.



Adding a Partition Scheme with File Groups to the Partition Function

```
USE AutoPartition
GO
CREATE PARTITION SCHEME PS MonthWise
AS PARTITION PF_MonthlyPartition
TO
 'FG_01_2020', 'FG_02_2020', 'FG_03_2020',
 'FG_04_2020', 'FG_05_2020', 'FG_06_2020',
 'FG_07_2020', 'FG_08_2020', 'FG_09_2020',
 'FG_10_2020', 'FG_11_2020', 'FG_12_2020',
 'Primary'
```

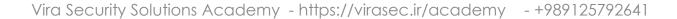


Adding a filegroup to the partition schema with attaching the partition function



Here, the Primary filegroup is an additional filegroup in the partition scheme definition.

The primary filegroup is used to store those rows which are exceeding the partition range in the function. It works when users forget to add new ranges and new filegroups with the file.



Creating a Table with the Partition Scheme

```
USE AutoPartition
GO
CREATE TABLE orders
 [order_id] BIGINT IDENTITY(1,1) NOT NULL
 [user id] BIGINT,
 [order_amt] DECIMAL(10,2),
 [address_id] BIGINT,
 [status_id] TINYINT,
 [is_active] BIT,
 [order date] [datetime]
) ON PS_MonthWise ([order_date]);
GO
CREATE CLUSTERED INDEX CI orders order id ON orders (order id)
GO
CREATE NONCLUSTERED INDEX IX user id ON orders(user id)
GO
```



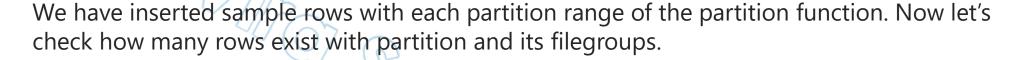
Here, the table is defined with applying the Partition to the column [order_date] of table orders.

Inserting data into Table [orders]

```
INSERT INTO orders (user id, order amt, address id, status id, is active, order date)
VALUES
(43, 1623.78, 51, 1, 1, '2020-01-14 13:21:51.869'),
(51, 8963.17, 43, 1, 1, '2020-02-17 05:07:43.193'),
(27, 7416.93, 66, 1, 1, '2020-03-21 07:53:07.743'),
(58, 9371.54, 45, 1, 1, '2020-04-26 16:19:27.852'),
(53, 8541.56, 65, 1, 1, '2020-05-08 19:21:58.654'),
(98, 6971.85, 54, 1, 1, '2020-06-17 21:34:52.426'),
(69, 5217.74, 78, 1, 1, '2020-07-03 07:37:51.391'),
(21, 9674.14, 98, 1, 1, '2020-08-27 23:49:53.813'),
(52, 1539.96, 32, 1, 1, '2020-09-01 17:17:07.317'),
(17, 7193.63, 21, 1, 1, '2020-10-23 10:23:37.307'),
(68, 3971.25, 19, 1, 1, '2020-11-30 09:01:27.079'),
(97, 5973.58, 97, 1, 1, '2020-12-06 13:43:21.190'),
(76, 4163.95, 76, 1, 1, '2021-01-03 18:51:17.764')
GO
```



Inserting data into Table [orders]





Partition details with Row count





```
SELECT DISTINCT o.name as table_name, rv.value as partition_range, fg.name as file_groupName, p.partition_num
ber, p.rows as number of rows
FROM sys.partitions p
INNER JOIN sys.indexes i ON p.object id = i.object id AND p.index id = i.index id
INNER JOIN sys.objects o ON p.object id = o.object id
INNER JOIN sys.system_internals_allocation_units au ON p.partition_id = au.container_id
INNER JOIN sys.partition schemes ps ON ps.data space id = i.data space id
INNER JOIN sys.partition_functions f ON f.function_id = ps.function_id
INNER JOIN sys.destination_data_spaces dds ON dds.partition_scheme_id = ps.data_space_id AND dds.destination
id = p.partition number
INNER JOIN sys.filegroups fg ON dds.data space id = fg.data space id
LEFT OUTER JOIN sys.partition range values rv ON f.function id = rv.function id AND p.partition number = rv.bou
ndary id
WHERE o.object id = OBJECT ID('orders');
```

Vinda Security Solid

Here, the partition range was not defined for January 2021. But we added the last file group with the Primary. Therefore, if rows are exceeding the range, then those will be allocated to the primary filegroup.

Table Rows with Partition Number

Users can find the partition number with each row of the table as well. Users can bifurcate the row allocation to the logical partition number with the help of a \$PARTITION() function.

Security Solution

SELECT \$PARTITION.PF_MonthlyPartition(order_date) AS PartitionNumber, * FROM orders

Automate the Partition flow

We have explained the above examples to understand the requirement of post activities and maintenance on the partition function. A primary filegroup is defined to manage those rows which are out of the partition range. However, the database team has to monitor that the range of any partition function is ending or not? To avoid some manual tasks, users can set up the SQL Server job to perform it automatically.

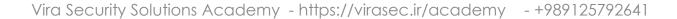
SQL Server job will be executed in a pre-defined scheduled time (monthly or weekly) and helps to find out the partition functions which are needed to be maintained. In the above sample, we used the DATETIME column type for the partition range. Users can write a program with T-SQL statements as below:

Find the Partition Function whose range does not exist for the next month. For example, in the above example, partition range is defined till Dec 2020, and a current timestamp is 2020-12-27 16:27:09.500. After three days, the partition range will be exceeded because the maximum range is 2020-12-31 23:59:59.997 for the order table. Now, we will find the partition functions which are required to be maintained using the below T-SQL.



Automate the Partition flow

```
SELECT o.name as table name,
 pf.name as PartitionFunction,
 ps.name as PartitionScheme,
 MAX(rv.value) AS LastPartitionRange,
 CASE WHEN MAX(rv.value) <= DATEADD(MONTH, 2, GETDATE()) THEN 1 else 0 END AS isRequiredMainten
ance
--INTO #temp
FROM sys.partitions p
INNER JOIN sys.indexes i ON p.object id = i.object id AND p.index id = i.index id
INNER JOIN sys.objects o ON p.object id = o.object id
INNER JOIN sys.system internals allocation units au ON p.partition id = au.container id
INNER JOIN sys.partition_schemes ps ON ps.data_space_id = i.data_space_id <
INNER JOIN sys.partition functions pf ON pf.function id = ps.function id
INNER JOIN sys.partition_range_values rv ON pf.function_id = rv.function_id AND p.partition_number = rv.
boundary id
GROUP BY o.name, pf.name, ps.name
```



Automate the Partition flow

Here, the above result set returned the partition function (**PF_MonthlyPartition**) for adding the new range.

The following code helps to insert information to the new temp table for those partition functions that are required to SPLIT.

```
SELECT table name,
 PartitionFunction,
 PartitionScheme,
 LastPartitionRange,
 CONVERT(VARCHAR, DATEADD(MONTH, 1, LastPartitionRange), 25) AS NewRange,
 'FG' + CAST(FORMAT(DATEADD(MONTH, 1, LastPartitionRange), MM') AS VARCHAR(2)) +
  CAST(YEAR(DATEADD(MONTH, 1, LastPartitionRange)) AS VARCHAR(4)) AS NewFileGroup,
 'File_'+ CAST(FORMAT(DATEADD(MONTH, 1, LastPartitionRange), 'MM') AS VARCHAR(2)) +
  CAST(YEAR(DATEADD(MONTH, 1, LastPartitionRange)) AS VARCHAR(4)) AS FileName,
 'C:\Program Files\Microsoft SQL Server\MSSQL13.JRAIYANI\MSSQL\DATA\' AS file path
INTO #generateScript
FROM #temp
WHERE isRequiredMaintenance = 1
```



We can also generate a dynamic script to create File Group, File, add a new file group to partition scheme, and new range to the partition function as below:





```
DECLARE @filegroup NVARCHAR(MAX) = "
DECLARE @file NVARCHAR(MAX) = "
DECLARE @PScheme NVARCHAR(MAX) = "
DECLARE @PFunction NVARCHAR(MAX) = "
SELECT @filegroup = @filegroup +
  CONCAT('IF NOT EXISTS(SELECT 1 FROM AutoPartition.sys.filegroups WHERE name = "',NewFileGroup,"')
  BEGIN
   ALTER DATABASE AutoPartition ADD FileGroup ', NewFileGroup,'
  END;'),
  @file = @file + CONCAT('IF NOT EXISTS(SELECT 1 FROM AutoPartition.sys.database files WHERE name = "',FileName,"')
  BEGIN
  ALTER DATABASE AutoPartition ADD FILE
  (NAME = "',FileName,"',
  FILENAME = "',File Path,FileName,'.ndf",
  SIZE = 5MB, MAXSIZE = UNLIMITED,
  FILEGROWTH = 10MB)
  TO FILEGROUP ', NewFileGroup, '
  END;'),
  @PScheme = @PScheme + CONCAT('ALTER PARTITION SCHEME', PartitionScheme, 'NEXT USED', NewFileGroup,';'),
  @PFunction = @PFunction + CONCAT('ALTER PARTITION FUNCTION', PartitionFunction, '() SPLIT RANGE ("', NewRange,
FROM #generateScript
EXEC (@filegroup)
EXEC (@file)
EXEC (@PScheme)
EXEC (@PFunction)
```





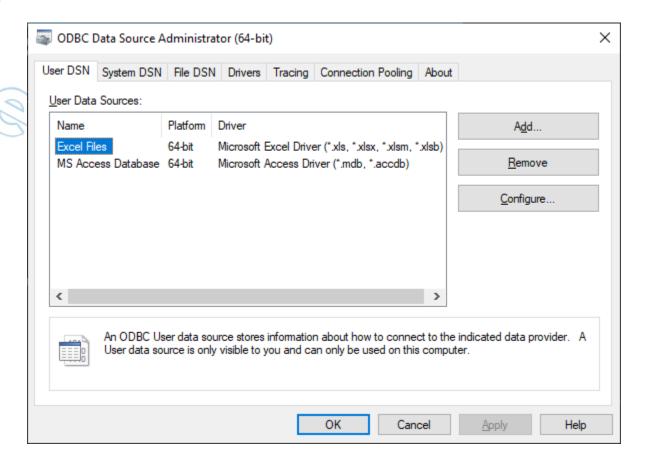
Security Solid

Telnet

ODBC connection Manager

sqlcmd

ODBC connection Manager





Sqlcmd

sqlcmd -S [ServerName]

sqlcmd -s nightingale

1>select name from sys.databases

